

CS306: Introduction to IT Security

Fall 2020

Lecture 6: MACs & Hashing

Instructor: **Nikos Triandopoulos**

October 6, 2020



6.0 Announcements

CS306: Other announcements

- ◆ HW2 to come by Friday this week
- ◆ Road ahead
 - ◆ no lecture on October 13 (next week, classes will run on Monday schedule)
 - ◆ regular lecture on October 20
 - ◆ midterm exam on October 27 (in whatever format)

CS306: Tentative Syllabus

Week	Date	Topics	Reading	Assignment
1	Sep 1	Introduction	Lecture 1	-
2	Sep 8	Symmetric-key encryption	Lecture 2	Lab 1
3	Sep 15	Perfect secrecy	Lecture 3	Lab 2, HW 1
4	Sep 22	Ciphers in practice I	Lecture 4	Lab 3, HW 1
5	Sep 29	Ciphers in practice II	Lecture 5	Lab 4
6	Oct 6	MACs & hashing		
-	Oct 13	No class (Monday schedule)		
7	Oct 20	Public-key cryptography		

CS306: Tentative Syllabus

(continued)

Week	Date	Topics	Reading	Assignment
8	Oct 27	Midterm	All materials covered	
9	Nov 3	Network/Web security		
10	Nov 10	Software/Database security		
11	Nov 17	Cloud security		
12	Nov 24	AC/Authentication/Privacy		
13	Dec 1	Economics		
14	Dec 8	Legal & ethical issues		
15	Dec 10 (or later)	Final (closed “books”)	All materials covered*	

Last week

- ◆ Ciphers in practice
 - ◆ Revision
 - ◆ the big picture, computational security, pseudo-randomness, stream ciphers, PRGs
 - ◆ Block ciphers, pseudorandom functions
 - ◆ Modes of operations
 - ◆ DES, AES
- ◆ Demo
 - ◆ The Caesar and Vigenère ciphers and their cryptanalysis (Afternoon)
 - ◆ Pseudo-randomness in practice (Evening)

Today

- ◆ Message authentication
 - ◆ MACs
 - ◆ Replay attacks
 - ◆ Constructions
- ◆ Cryptographic hashing
 - ◆ Hash functions
 - ◆ Constructions
- ◆ Demo
 - ◆ Hash functions in practice

6.1 Message authentication

Recall: Integrity

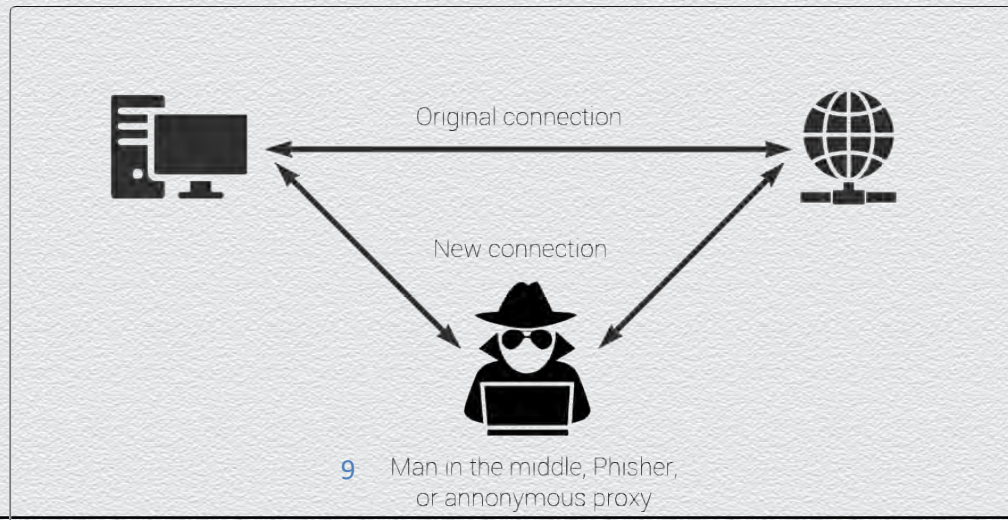
Fundamental security property

- ◆ **an asset is modified only by authorized parties**
- ◆ “I” in the CIA triad

*“computer security seeks to prevent **unauthorized** viewing (confidentiality) or **modification (integrity)** of **data** while preserving access (availability)”*

Alteration

- ◆ main threat against integrity of **in-transit** data
- ◆ e.g., MITM attack



Security problems studied by modern cryptography

- ◆ Classical cryptography: **message encryption**
 - ◆ early crypto schemes tried to provide **secrecy / confidentiality**
- ◆ Modern cryptography: **wide variety** of security problems
 - ◆ today we need to study a large set of **security properties** beyond secrecy
- ◆ The sibling of message encryption: **message authentication**
 - ◆ another cornerstone of any secure system aiming to provide **authenticity & integrity**

Message authentication: Motivation

Information has **value**, but only when it is **correct**

- ◆ random, incorrect, inaccurate or maliciously altered data is **useless** or **harmful**
 - ◆ **message authentication = message integrity + authenticity**
 - ◆ while in transit (or at rest), no message should be **modified** by an outsider
 - ◆ no outsider can **impersonate** the stated message sender (or owner)
- ◆ it is often necessary / worth to protect critical / valuable data
 - ◆ **message encryption**
 - ◆ while in transit (or at rest), no message should be **leaked** to an outsider

Example 1

Secure electronic banking

- ◆ a bank receives an electronic request to transfer \$1,000 from Alice to Bob

Concerns

- ◆ who ordered the transfer, Alice or an attacker (e.g., Bob)?
- ◆ is the amount the intended one or was maliciously modified while in transit?
 - ◆ adversarial Vs. random message-transmission errors
 - ◆ standard error-correction is **not sufficient** to address this concern

Example 2

Web browser cookies

- ◆ a user is performing an online purchase at Amazon
- ◆ a “cookie” contains session-related info, as client-server HTTP traffic is stateless
 - ◆ stored at the client, included in messages sent to server
 - ◆ contains client-specific info that affects the transaction
 - ◆ e.g., the user’s shopping cart along with a discount due to a coupon

Concern

- ◆ was such state maliciously altered by the client (possibly harming the server)?

Integrity of communications / computations

Highly important

- ◆ any unprotected system cannot be assumed to be trustworthy w.r.t.
 - ◆ origin/source of information (due to impersonation attacks, phishing, etc.)
 - ◆ contents of information (due to man-in-the-middle attacks, email spam, etc.)
 - ◆ overall system functionality

Prevention Vs. detection

- ◆ unless system is “closed,” adversarial tampering with its integrity **cannot be avoided!**
- ◆ goal: identify system components that are not trustworthy
 - ◆ **detect tampering** or **prevent undetected tampering**
 - ◆ e.g., avoid “consuming” falsified information

Encryption does not imply authentication

A common misconception

“since ciphertext c hides message m , Mallory cannot meaningfully modify m via c ”

Why is this incorrect?

- ◆ all encryption schemes (seen so far) are based on one-time pad, i.e., masking via XOR
- ◆ consider flipping a single bit of ciphertext c ; what happens to plaintext m ?
 - ◆ such property of one-time pad does not contradict the secrecy definitions

Generally, secrecy and integrity are distinct properties

- ◆ encrypted traffic generally provides **no integrity** guarantees

6.2 Message authentication codes (MACs)


Problem setting: Reliable communication

Two parties wish to communicate over a channel

- ◆ Alice (sender/source) wants to send a message m to Bob (recipient/destination)

Underlying channel is unprotected


- ◆ Mallory (attacker/adversary) can manipulate any sent messages
- ◆ e.g., message transmission via a compromised router



Solution concept: Symmetric-key message authentication

Main idea

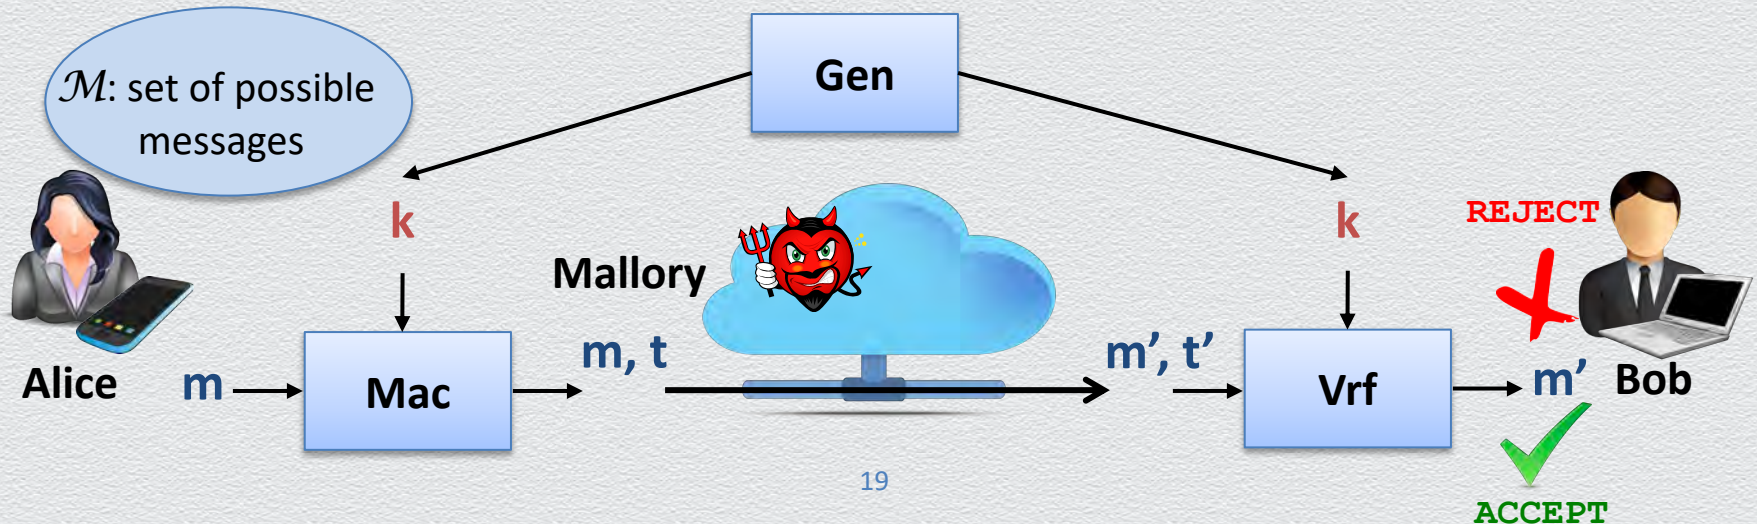
- ◆ secretly annotate or “sign” message so that it is **unforgeable** while in transit
 - ◆ Alice **tags** her message m with **tag** t , which is sent **along** with **plaintext** m
 - ◆ Bob **verifies** authenticity of received message using tag t
 - ◆ Mallory can manipulate m, t but “**cannot forge**” a fake verifiable pair m', t'
 - ◆ Alice and Bob share a **secret key** k that is used for both operations



Security tool: Symmetric Message Authentication Code

Abstract cryptographic primitive, a.k.a. **MAC**, defined by

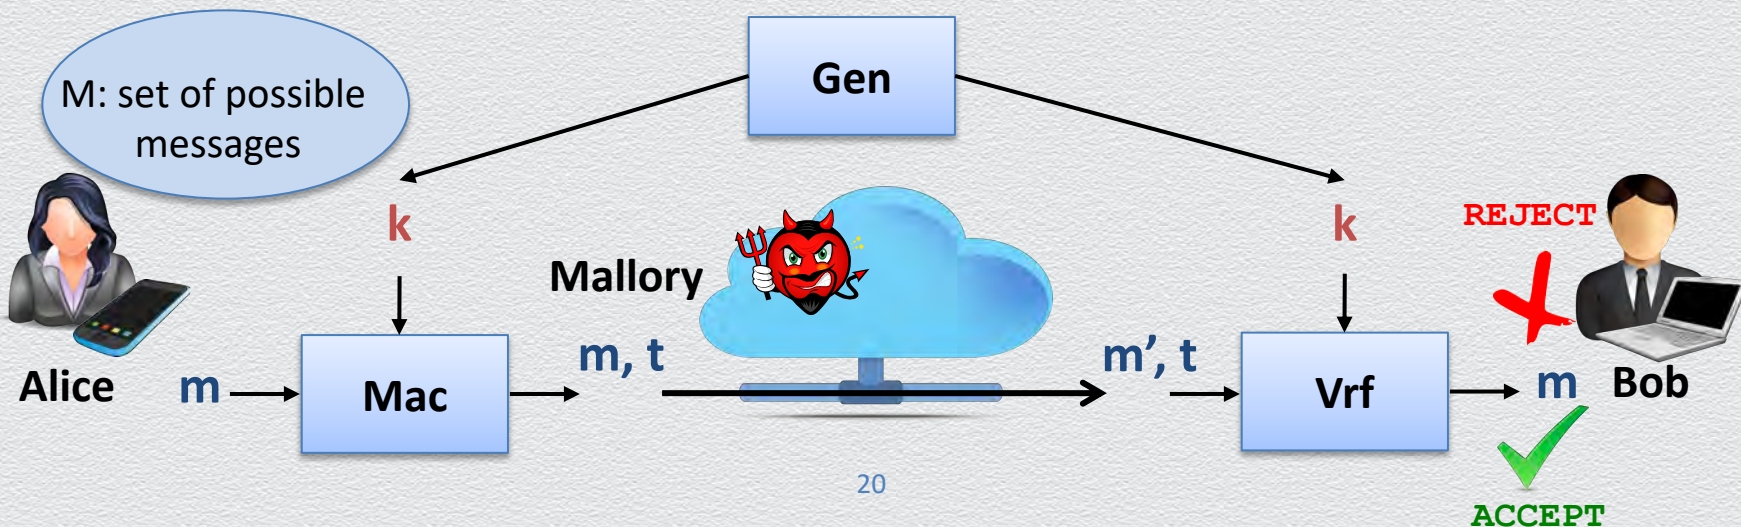
- ◆ a **message space** \mathcal{M} ; and
- ◆ a triplet of algorithms (**Gen**, **Mac**, **Vrf**)
 - ◆ Gen, Mac are probabilistic algorithms, whereas Vrf is deterministic
 - ◆ Gen outputs a uniformly random key k (from some key space \mathcal{K})



Desired properties for MACs

By design, any MAC should satisfy the following

- ◆ **efficiency:** key generation & message transformations “are fast”
- ◆ **correctness:** for all m and k , it holds that $\text{Vrf}_k(m, \text{Mac}_k(m)) = \text{ACCEPT}$
- ◆ **security:** one “cannot forge” a fake verifiable pair m', t'



Main application areas


Secure communication

- ◆ **verify authenticity of messages** sent among parties
- ◆ assumption
 - ◆ Alice and Bob **securely generate, distribute and store shared key k**
 - ◆ attacker does not learn key k



Secure storage

- ◆ **verify authenticity of files** outsourced to the cloud
- ◆ assumption
 - ◆ Alice **securely generates and stores key k**
 - ◆ attacker does not learn key k



Conventions

Random key selection

- ◆ typically, Gen selects key k **uniformly at random** from the key space \mathcal{K}

Canonical verification

- ◆ when Mac is deterministic, Vrf typically amounts to re-computing the tag t
 - ◆ $\text{Vrf}_k(m, t)$: 1. $t' := \text{Mac}_k(m)$ 2. if $t = t'$, output `ACCEPT` else output `REJECT`
- ◆ but conceptually the following operations are distinct
 - ◆ authenticating m (i.e., running Mac) Vs. verifying authenticity of m (i.e., running Vrf)

MAC security

MAC scheme
(Gen, Mac, Vrf)



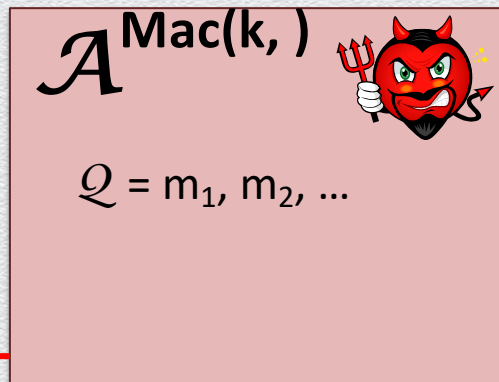

\mathcal{T}

Gen \rightarrow k

Mac_k(m_i) \rightarrow t_i

Attacker **wins** the game if

1. Vrf_k(m*, t*) = ACCEPT &
2. m* not in \mathcal{Q}



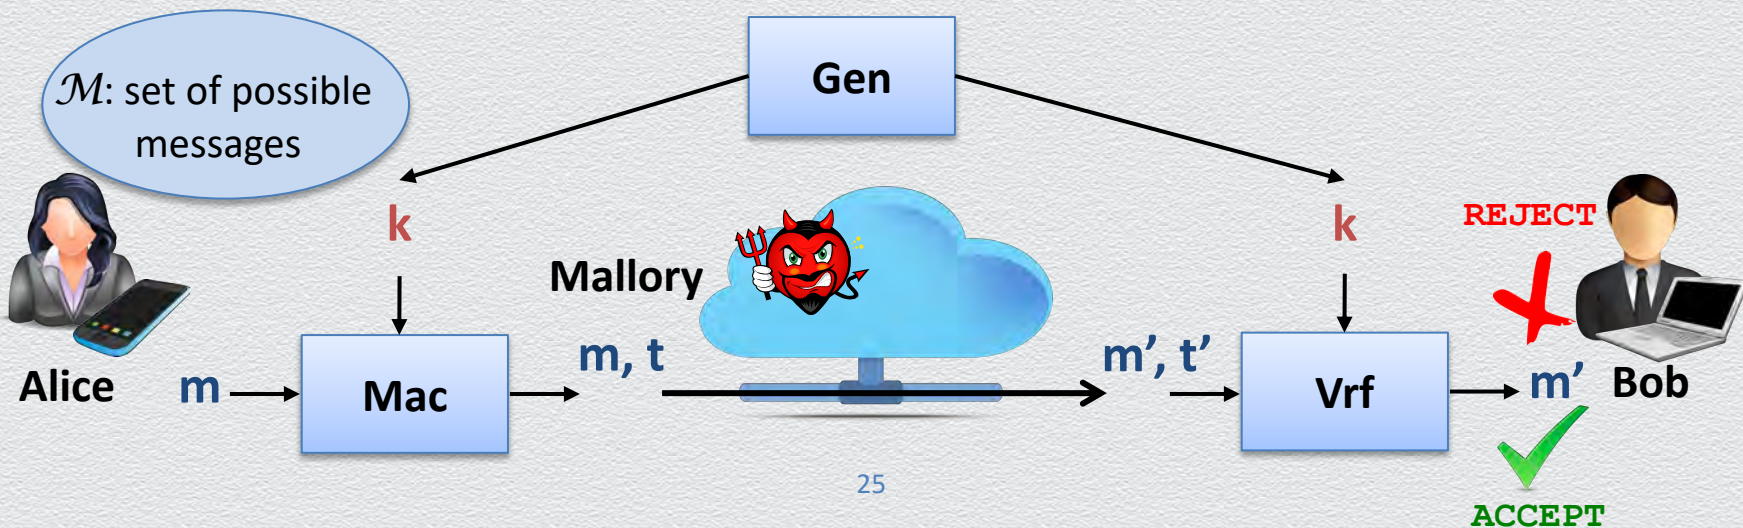
The MAC scheme is **secure** if any PPT \mathcal{A} wins the game only negligibly often.

6.3 Replay attacks

Recall: MAC

Abstract cryptographic primitive, a.k.a. **MAC**, defined by

- ◆ a **message space** \mathcal{M} ; and
- ◆ a triplet of algorithms (**Gen**, **Mac**, **Vrf**)



Recall: MAC security

MAC scheme
(Gen, Mac, Vrf)



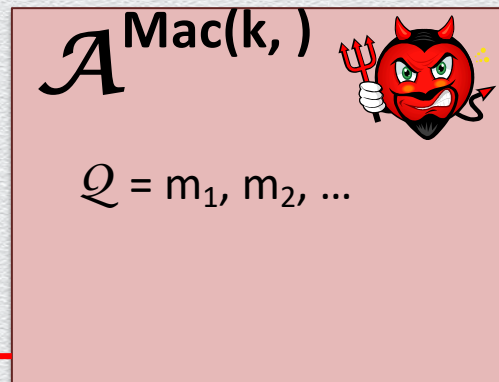

\mathcal{T}

$\text{Gen} \rightarrow k$

$\text{Mac}_k(m_i) \rightarrow t_i$

Attacker **wins** the game if

1. $\text{Vrf}_k(m^*, t^*) = \text{ACCEPT}$ &
2. m^* not in \mathcal{Q}



The MAC scheme is **secure** if any PPT \mathcal{A} wins the game only negligibly often.

Real-life attacker

In practice, an attacker may

- ◆ observe a traffic of authenticated (and successfully verified) messages
- ◆ manipulate (or often also partially influences) traffic
 - ◆ aims at inserting an invalid but verifiable message m^*, t^* into the traffic
 - ◆ interesting case: forged message is a new (unseen) one
 - ◆ trivial case: forged message is a previously observed one, a.k.a. a **replay attack**
- ◆ launch a **brute-force attack** (given that $\text{Mac}_k(m) \rightarrow t$ is publicly known)
 - ◆ given any observed pair m, t , exhaustively search key space to find the used key k

Threat model

In the security game, Mallory is an adversary \mathcal{A} who is

- ◆ “active” (on the wire)
 - ◆ we allow \mathcal{A} to **observe** and **manipulate** sent messages
- ◆ “well-informed”
 - ◆ we allow \mathcal{A} to **request MAC tags** of messages of **its choice**
- ◆ “replay-attack safe”
 - ◆ we restrict \mathcal{A} to **forge only new** messages
- ◆ “PPT”
 - ◆ we restrict \mathcal{A} to be **computationally bounded**
 - ◆ new messages may be forged undetectably only negligibly often

Notes on security definition

Is it a rather strong security definition?

- ◆ we allow \mathcal{A} to **query MAC tags for any message**
 - ◆ but real-world senders will authenticate only “meaningful” messages
- ◆ we allow \mathcal{A} to break the scheme by **forging any new message**
 - ◆ but real-world attackers will forge only “meaningful” messages

Yes, it is the right approach...

- ◆ message “meaningfulness” **depends on higher-level application**
 - ◆ text messaging apps require authentication of English-text messages
 - ◆ other apps may require authentication of binary files
 - ◆ security definition should better be **agnostic** of the specific higher application

Notes on security definition (II)

Are replay attacks important in practice?

- ◆ absolutely yes: a **very realistic & serious threat!**
 - ◆ e.g., what if a money transfer order is “replayed”?

Yet, a “replay-attack safe” security definition is preferable

- ◆ again, whether replayed messages are valid depends on higher-level app
- ◆ better to delegate to this app the specification of such details
 - ◆ e.g., semantics on traffic or validity checks on messages before they’re “consumed”

Eliminating replay attacks

- ◆ use of counters (i.e., common shared state) between sender & receiver
- ◆ use of timestamps along with a (relaxed) authentication window for validation


6.4 MAC constructions

Three generic MAC constructions

- ◆ fixed-length MAC
 - ◆ direct application of a PRF for tagging
 - ◆ limited applicability
- ◆ domain extension for MACs
 - ◆ straightforward secure extension of fix-length MAC
 - ◆ inefficient
- ◆ CBC-MAC
 - ◆ resembles CBC-mode encryption
 - ◆ efficient

1. Fixed-length MAC

- ◆ based on use of a PRF
 - ◆ employ a PRF F_k in the obvious way to compute and canonically verify tags
 - ◆ set tag t to be the pseudorandom string derived by evaluating F_k on message m
- ◆ secure, provided that F_k is a secure PRF



MAC scheme Π

$\text{Gen}(1^n): \{0,1\}^n \rightarrow k$

$\text{Mac}_k(m): \text{set } t = F_k(m)$


$\text{Vrfy}_k(m,t): \text{return } 1 \text{ iff } t = F_k(m)$

2. Domain extension for MACs (I)

- ◆ suppose we have the previous fix-length MAC scheme
- ◆ how can we authenticate a message m of arbitrary length?

- ◆ naïve approach

- ◆ pad m and view it as d blocks m_1, m_2, \dots, m_d
- ◆ separately apply MAC to block m_i



- ◆ security issues

- ◆ reordering attack; verify block index, $t = F_k(m_i || i)$
- ◆ truncation attack; verify message length $\delta = |m|$, $t = F_k(m_i || i || \delta)$
- ◆ mix-and-match attack; randomize tags (using message-specific fresh nonce)


2. Domain extension for MACs (II)

Final scheme

- ◆ assumes a secure MAC scheme for messages of size n
- ◆ set tag of message m of size δ at most $2^{n/4}$ as follows
 - ◆ choose fresh random nonce r of size $n/4$; view m as d blocks of size $n/4$ each
 - ◆ separately apply MAC on each block, authenticating also its index, δ and nonce r

Security

- ◆ extension is secure, if F_k is a secure PRF




3. CBC-MAC

Idea

- ◆ employ a PRF in a manner similar to CBC-mode encryption


Security

- ◆ extension is secure, if
 - ◆ F_k is a secure PRF; and
 - ◆ only **fixed-length** messages are authenticated
- ◆ messages of length equal to any multiple of n can be authenticated
 - ◆ but this length need be fixed in advance
 - ◆ insecure, otherwise




3. CBC-MAC Vs. previous schemes

- ◆ can authenticate longer messages than basic PRF-based scheme (1)




- ◆ more efficient than domain-extension MAC scheme (2)

Scheme (1)




Scheme (2)




3. CBC-MAC Vs. CBC-mode encryption

- ◆ crucially for their security
 - ◆ CBC-MAC uses **no IV** (or uses an IV set to 0) and only the **last PRF output**
 - ◆ CBC-mode encryption uses a **random IV** and **all PRF outputs**
 - ◆ “simple”, innocent modification can be catastrophic...

CBC-MAC



CBC-mode encryption




6.5 Hash functions

Cryptographic hash functions

Basic cryptographic primitive

- ◆ maps “**objects**” to a **fixed-length** binary strings
- ◆ core security property: mapping **avoids collisions**
 - ◆ **collision**: distinct objects ($x \neq y$) are mapped to the same hash value ($H(x) = H(y)$)
 - ◆ although collisions **necessarily exist**, they are **infeasible to find**




Important role in modern cryptography

- ◆ lie between symmetric- and asymmetric-key cryptography
- ◆ capture different security properties of “idealized random functions”
- ◆ qualitative stronger assumption than PRF

Hash & compression functions


Map messages to short digests

- ◆ a **general** hash function $H()$ maps
 - ◆ a message of an arbitrary length to a $l(n)$ -bit string
- ◆ a **compression** (hash) function $h()$ maps
 - ◆ a long binary string to a shorter binary string
 - ◆ an $l'(n)$ -bit string to a $l(n)$ -bit string, with $l'(n) > l(n)$



Collision resistance (CR)

Attacker wins the game if $x \neq x' \text{ \& } H(x) = H(x')$



H is collision-resistant if any PPT \mathcal{A} wins the game only negligibly often.

Weaker security notions

Given a hash function $H: X \rightarrow Y$, then we say that H is


- ◆ **preimage resistant** (or **one-way**)
 - ◆ if given $y \in Y$, finding a value $x \in X$ s.t. $H(x) = y$ happens negligibly often
- ◆ **2-nd preimage resistant** (or **weak collision resistant**)
 - ◆ if given a uniform $x \in X$, finding a value $x' \in X$, s.t. $x' \neq x$ and $H(x') = H(x)$ happens negligibly often
- ◆ cf. **collision resistant** (or **strong collision resistant**)
 - ◆ if finding two distinct values $x', x \in X$, s.t. $H(x') = H(x)$ happens negligibly often

6.6 Design framework

Domain extension via the Merkle-Damgård transform

General design pattern for cryptographic hash functions

- ◆ reduces CR of general hash functions to CR of compression functions



- ◆ thus, in practice, it suffices to realize a collision-resistant compression function h
- ◆ compressing by 1 single bit is at least as hard as compressing by any number of bits!


Merkle-Damgård transform: Design

Suppose that $h: \{0,1\}^{2n} \rightarrow \{0,1\}^n$ is a collision-resistant compression function

Consider the general hash function $H: \mathcal{M} = \{x : |x| < 2^n\} \rightarrow \{0,1\}^n$, defined as

Merkle-Damgård design


- ◆ $H(x)$ is computed by applying $h()$ in a **“chained” manner** over n -bit message blocks



- ◆ pad x to define a number, say B , **message blocks x_1, \dots, x_B** , with $|x_i| = n$
- ◆ set extra, final, message block **x_{B+1} as an n -bit encoding L of $|x|$**
- ◆ starting by initial digest **$z_0 = IV = 0^n$** , output **$H(x) = z_{B+1}$** , where **$z_i = h^s(z_{i-1} || x_i)$**

Merkle-Damgård transform: Security

If the compression function h is CR,
then the derived hash function H is also CR!



Compression function design: The Davies-Meyer scheme


Employs PRF w/ key length m & block length n

- ◆ define $h: \{0,1\}^{n+m} \rightarrow \{0,1\}^n$ as

$$h(x \parallel k) = F_k(x) \text{ XOR } x$$

Security

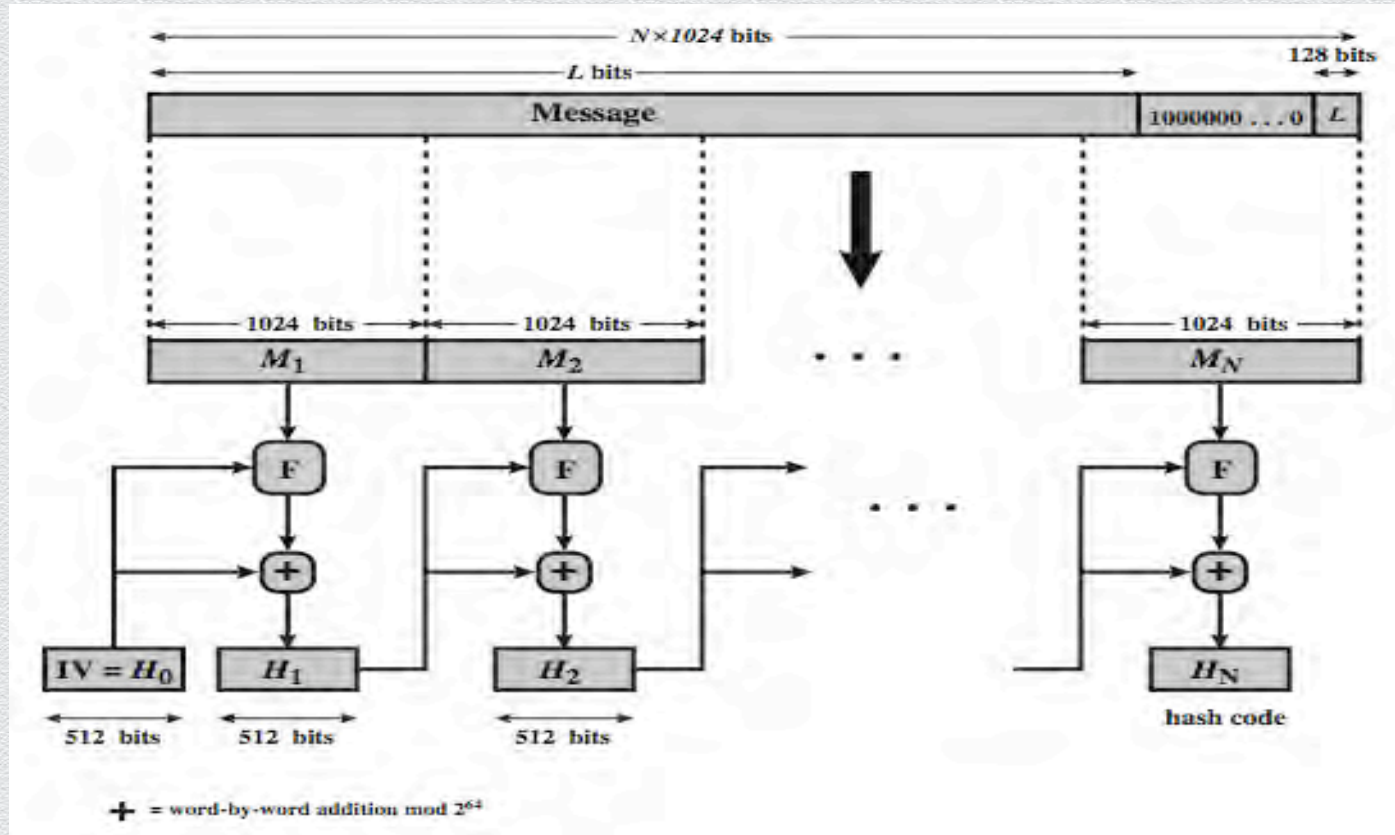
- ◆ h is CR, if F is an **ideal cipher**



Well known hash functions

- ◆ MD5 (designed in 1991)
 - ◆ output 128 bits, collision resistance **completely broken** by researchers in 2004
 - ◆ today (controlled) collisions can be found in less than a minute on a desktop PC
- ◆ SHA1 – the Secure Hash Algorithm (series of algorithms standardized by NIST)
 - ◆ output 160 bits, considered **insecure** for collision resistance
 - ◆ **broken** in 2017 by researchers at CWI
- ◆ SHA2 (SHA-224, SHA-256, SHA-384, SHA-512)
 - ◆ outputs 224, 256, 384, and 512 bits, respectively, **no real security concerns yet**
 - ◆ based on Merkle-Damgård + Davies-Meyer generic transforms
- ◆ SHA3 (Kessac)
 - ◆ **completely new philosophy** (sponge construction + unkeyed permutations)

SHA-2-512 overview



Current hash standards

Algorithm	Maximum Message Size (bits)	Block Size (bits)	Rounds	Message Digest Size (bits)
MD5	2^{64}	512	64	128
SHA-1	2^{64}	512	80	160
SHA-2-224	2^{64}	512	64	224
SHA-2-256	2^{64}	512	64	256
SHA-2-384	2^{128}	1024	80	384
SHA-2-512	2^{128}	1024	80	512
SHA-3-256	unlimited	1088	24	256
SHA-3-512	unlimited	576	24	512